

# Neural Network Basics

José Dorronsoro  
Escuela Politécnica Superior  
Universidad Autónoma de Madrid

2021

## Contents

<b>1</b>	<b>Multilayer Perceptrons</b>	<b>3</b>
1.1	Basic Modeling . . . . .	3
1.2	Regression Basics . . . . .	6
1.3	Cross Validation, Bias and Variance . . . . .	13
1.4	Basic Classification . . . . .	17
<b>2</b>	<b>Multilayer Perceptrons</b>	<b>24</b>
2.1	Rosenblatt's Perceptrons . . . . .	24
2.2	Classical MLPs . . . . .	32
2.3	MLP Regularization . . . . .	44
2.4	Unconstrained Smooth Optimization . . . . .	46
2.5	Computational Costs of MLPs . . . . .	52

# 1 Multilayer Perceptrons

## 1.1 Basic Modeling

### What Is Machine Learning (ML)?

- Lofty definition: make machines learn!!!
  - Have to make “machines” and “learn” more precise
- The machines of ML: mathematical input–output processes that lend themselves to some form of (numerical) parameterization
- The learning process: adjust the machine’s parameters until a goal is reached
- New thing: “goal”?
  - At first sight, get something done
  - Ultimately, to minimize some error measure
- Summing things up: a ML process tries to find a concrete mathematical/algorithmic **input–output parameterized transformation** that **minimizes an error measure** by iteratively **adjusting the transformation’s parameters**

### Where Lies ML?

- In the middle of a possibly long process chain
- Before ML starts we must
  - Go from **raw to organized** data: accessing, gathering, cleaning, formatting, ...
  - Go from **organized to** (potentially) **informative** data: extracting basic and derived features
- After ML finishes and start applying a model, we must perform
  - Outcome **evaluation**: how good/actionable is
  - Outcome **exploitation**: collect, organize, act
  - **Individual model maintenance**: monitor performance, tune hyper–parameters
  - **Modeling life cycle maintenance**: discard old models, introduce new ones and **communicate** our work/results

### Supervised/Unsupervised Models

- ML model types: **supervised, unsupervised**
- Supervised models:
  - Targets  $y^p$  are known and the model tries to predict or estimate them
  - These known targets guide, or **supervise**, model building
  - Main emphasis here

- Unsupervised models:
  - There are no predetermined or supervising outputs
  - But nevertheless the model is supposed to learn relations or find structure in the data
  - Often as a first step towards a supervised model
- Semisupervised models lie in between: some patterns have targets, some don't

### Regression and Classification

- Problems (usually) to be solved by ML models: regression, classification
- In both patterns come in pairs  $(x, y)$ 
  - $x$ : inputs, predictors, features, independent variables
  - $y$ : target, response, dependent variable; numerical in regression, class labels in classification
- **Regression**: the desired output  $y$  is **regressed** into the inputs  $x$  to derive a model  $\hat{y} = f(x)$ 
  - We want  $y \simeq \hat{y}$  so having  $y - \hat{y}$  “small” is the natural goal
- **Classification**: inputs are derived from several classes  $C_1, \dots, C_K$ , to which labels  $\ell_k$  are assigned
  - The model now assigns a label  $\ell(x)$  to an input  $x$
  - If  $x$  is derived from  $C_k$  we want to have  $\ell(x) = \ell_k$
  - Here having  $\ell(x) - \ell_k$  “small” may not make sense


### The Boston Housing Regression Problem

- This is a first, widely used “toy” problem
- We want to estimate the median of house values over an area from some information about it which we believe relevant
- Features  $x$ : several real estate-related variables of Boston areas
  - CRIM: per capita crime rate by town
  - RM: average number of rooms per dwelling
  - NOX: nitric oxides concentration (parts per 10 million)
  - AGE: proportion of owner-occupied units built prior to 1940
  - LSTAT: % lower status of the population
  - ...
- Target  $y$ : MEDV, median value of owner-occupied homes in \$1,000's

### The MNIST Classification Problem

- This is a well known benchmark, moderately large classification problem

- We have  $28 \times 28$  grey scale image rasters with handwritten digits 0 to 9
- We want to build a classifier that assigns each raster to one of the 0 to 9 digits
- Features  $x$ :  $28 \times 28$  image matrices
- Target  $y$ : labels  $[0, 1, \dots, 9]$
- Examples (0-padded to  $32 \times 32$ ):



### Model Parameterization

- Often individual models are selected through (ideally optimal) **parameter sets**
  - The parameters (weights)  $w \in R^M$  select a concrete  $f$  in a model family  $\mathcal{F}$
- **Parametric** models have a fixed functional form  $f(x) = f(x; W)$ 
  - Simplest example: **linear regression**, where  $x$  has dimension  $d$ ,  $M = d + 1$  and  $W = (w_0, w)$

$$f(x; w_0, w) = w_0 + \sum_{j=1}^d w_j x_j = w_0 + w \cdot x$$

- **Semi-parametric** models: also use weights but without a predefined functional form
  - MLPs but also RF or GBR
- **Non parametric** models do not use weights nor follow any broad functional form; Nearest Neighbor models

### Issues in Model Building

- There are some initial questions when working with models from a given family  $\mathcal{F}$ :
  - Which assumptions do they make?
  - How do they work on the input features?
  - How we do build them?
- In turn, these questions lead to another two:
  - How do we select the best model from the given family for the problem at hand?
  - How do we control the model building procedure?
- All of them address fundamental issues that require a **moderately deep** understanding of what is going on under the model's hood
- This understanding is usually framed in **mathematical language**

## 1.2 Regression Basics

### How to Build Regression Models

- In general we have a sample  $S = \{x^p, y^p\}, 1 \leq p \leq N$ , with  $x^p$  the **features** and  $y^p$  the **targets**
- We want to build a model  $\hat{y} = f(x)$  so that  $\hat{y}^p = f(x^p) \simeq y^p$ ;
  - I.e., we want to **regress**  $y$  to the  $f(x)$
- The concrete  $f$  is chosen within a certain family  $\mathcal{F}$ 
  - Examples here: linear regression, multilayer perceptrons (MLPs)
  - And also: SVMs, Random Forests (RF), Gradient Boosting (GB), nearest neighbor (NN)
- Natural option to ensure  $f(x^p) \simeq y^p$ : choose  $f$  to minimize the sample **Mean Square Error (MSE)**

$$\hat{e}(f) = \hat{e}_S(f) = \frac{1}{2N} \sum_{p=1}^N (y^p - f(x^p))^2$$

- Thus, the model we select is  $\hat{f} = \hat{f}_S = \arg \min_{f \in \mathcal{F}} \hat{e}_S(f)$

### Model Estimation as Error Minimization

- For a parametric or semiparametric  $f(x; W)$  we can write  $\hat{e}_S(f) = \hat{e}_S(f(\cdot; W)) = \hat{e}_S(W)$
- The problem to solve becomes

$$\widehat{W}^* = \widehat{W}_S^* = \arg \min_W \hat{e}_S(W), \text{ i.e., } \hat{e}_S(\widehat{W}^*) \leq \hat{e}_S(W) \forall W$$

- In linear regression (LR) the sample error is

$$\hat{e}(w_0, w) = \frac{1}{2N} \sum_p (y^p - w_0 - w \cdot x^p)^2$$

i.e., the **sample** error, with a corresponding **population** error

$$e(w_0, w) = \frac{1}{2} E_{x,y} [(y - f(x))^2] = \frac{1}{2} E [(y - w_0 - w \cdot x)^2 p(x, y) dx dy]$$

- LR reduces to **minimize the quadratic form**  $\hat{e}_S(W)$

### Regression Assumptions

- **Key assumption:**  $x$  and  $y$  are related as  $y = \phi(x) + n$  where
  - $\phi(x)$  is the **true** underlying function
  - $n$  is **additive noise** with 0 mean and finite variance  $\sigma_N^2$
- Our sample is just a particular instance of a deeper **sample generation process**
- Thus  $x, n$  are produced by **random variables**  $X, N$

- And so is  $y$ , given by  $Y = \phi(X) + N$
- Moreover,  $X$  and  $N$  are **independent distributions** with densities  $q(x)$ ,  $\nu(n)$
- Thus,  $X$  and  $Y$  (or  $X$  and  $N$ ) have a joint density

$$p(x, y) = p(x, \phi(x) + n) = q(x) \nu(n) = q(x) \nu(y - \phi(x))$$

### MSE Decomposition

- We can decompose the MSE error of any model  $f$  as

$$\begin{aligned} 2\text{mse}(f) &= E_{x,y}[(y - f(x))^2] = \int (n + \phi(x) - f(x))^2 q(x) \nu(n) dx dn \\ &= \int (n^2 + 2n(\phi(x) - f(x)) + (\phi(x) - f(x))^2) q(x) \nu(n) dx dn \\ &= \int n^2 \nu(n) dn + \int (\phi(x) - f(x))^2 q(x) dx + \\ &\quad 2 \left( \int n \nu(n) dn \right) \left( \int (\phi(x) - f(x)) q(x) dx \right) \\ &= \sigma_N^2 + E_x[(\phi(x) - f(x))^2] \end{aligned}$$

- Thus **for any model we have  $\text{mse}(f) \geq \sigma_N^2$  always**
- And we should focus on achieving on  $f \simeq \phi$  (which we don't know!!)

### The Best Regression Model

- It is easy to see that the best  $f$  is simply  $f(x) = E_y[y|x]$ , for


$$E_y[y|x] = E_n[\phi(x) + n] = \int (\phi(x) + n) \nu(n) dn = \phi(x)$$

- Have we finished? In theory yes; in practice, not at all!!!
  - We do not know  $\nu$  and, thus, cannot compute the required integral
  - If we would have several  $M$  values  $y^j$  for any  $x$ , we could try  $\hat{\phi}(x) = \frac{1}{M} \sum_1^M y^j$
  - But this doesn't happen either
- Now we have two options:
  - Try to stretch the  $E[y|x]$  approach, for instance, by  **$k$ -NN Regression**
  - Forget about it and get back to get models  $f$  such that  $f \simeq \phi$

### From The Best Regressor to $k$ -NN

- The left hand side shows the ideal situation but the right one is what we should expect

### $k$ -NN Regression



- We can try to use the  $y^p$  values of several  $x^p$  close to a new  $x$
- We fix  $k$  neighbors  $x^{p_1}, \dots, x^{p_k}$  of  $x$  and estimate  $\hat{y} = \hat{y}(x)$  as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- $\hat{y}(x) = \hat{Y}_k^{NN}(x)$  is the  $k$ -**Nearest Neighbor (NN)** regressor
  - This can be refined to weighted versions, such as

$$\hat{y}(x) = \frac{1}{C_k(x)} \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2} y^{p_j}$$

with  $C_k(x) = \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2}$  a normalizing constant

- But we need that close predictors give predictions that are also close, **and** that there are enough of them close by
  - This is **very unlikely**

### The Curse of Dimensionality

- Even for low dimensions and large samples, **the sample space is essentially empty**



- Assume we have 1,000  $d$ -dimensional  $x$  patterns whose features have values between 1 and 10
  - In dimension  $d = 1$  there are 100 patterns per unit length
  - But when  $d = 3$  we have just 1 pattern per volume unit
  - And if  $d = 6$  we have just 1 pattern per 1,000 units of volume
  - And in dimension 10 (not a big one nowadays) we have just ... !!!
- Thus, for most problems, **there never will be enough close points**
- As a consequence, to get  $k$  observations we may go too far away from  $x$  and the average will not be meaningful
- Therefore, unless we deal with violently non-linear problems, a simple linear model may be better than  $k$ -NN regression for moderate dimensions

### Linear Models

- Assuming  $x \in R^d$ , the basic linear model is

$$f(x) = w_0 + \sum_1^d w_i x_i = w_0 + w \cdot x$$

- $w_0$  complicates notation; to drop it we center  $x$  and  $y$  so that  $E[x_i] = E[y] = 0$ ; then  $w_0 = 0$
- Then we are left with the simpler **homogeneous** model  $f(x) = w \cdot x$
- In practice we will always **normalize**  $x$ , for instance to have 0 mean and 1 standard deviation (std) on each feature
  - But not  $y$  if we may help it (or it is easy to reverse this)
- But: how do we find  $w$ ?

### 1-dimensional Linear Regression (LR)

- Assume that features  $X$  and target  $Y$  are **centered**, i.e., have 0 means
- For 1-dimensional patterns  $x$  the LR model then becomes  $f(x) = w x$
- And the error is then the function

$$\begin{aligned} \hat{e}(w) &= \frac{1}{2N} \sum_{p=1}^N (w x^p - y^p)^2 = \frac{1}{2N} \sum_p (w^2 (x^p)^2 - 2x^p y^p w + (y^p)^2) \\ &= w^2 \left( \frac{1}{2N} \sum_p (x^p)^2 \right) - w \left( \frac{1}{N} \sum_p x^p y^p \right) + \frac{1}{2N} \sum_p (y^p)^2 \end{aligned}$$

- Thus,  $\hat{e}(w) = aw^2 + bw + c$  with  $a > 0$  and it has obviously a minimum  $w^*$
- To find it we just solve  $\hat{e}'(w) = 0$

**Solving**  $\hat{e}'(w) = 0$

- We first compute  $\hat{e}'(w)$ , for which we have

$$\hat{e}'(w) = w \left( \frac{1}{N} \sum_p (x^p)^2 \right) - \frac{1}{N} \sum_p x^p y^p$$

- The optimal  $w^*$  solves  $\hat{e}'(w) = 0$  and is given by

$$w^* = \frac{\frac{1}{N} \sum_p x^p y^p}{\frac{1}{N} \sum_p (x^p)^2} = \frac{\frac{1}{N} X^t Y}{\frac{1}{N} X^t X} = \frac{\text{covar}(x, y)}{\text{var}(x)} = (\text{var}(x))^{-1} \text{covar}(x, y)$$

where  $X$  and  $Y$  denote the  $N \times d$  **data matrix** (vector here) and the  $N \times 1$  **target vector**

### General Linear Regression

- Assume again that  $X$  and  $Y$  are centered
- The LR model becomes now  $f(x) = \sum_1^d w_i x_i = w \cdot x$
- If  $Y$  is the  $N \times 1$  **target** vector and we organize the sample  $S$  in a  $N \times d$  **data matrix**  $X$ , the sample mse is given by

$$\begin{aligned} \hat{e}(w) &= \frac{1}{2N} \sum_p (w \cdot x^p - y^p)^2 = \frac{1}{2N} (Xw - Y)^t (Xw - Y) \\ &= \frac{1}{2N} (w^t X^t X w - 2w^t X^t Y + Y^t Y) \end{aligned}$$

- Now we have to solve  $\nabla \hat{e}(w) = 0$ , i.e.,  $\frac{\partial \hat{e}}{\partial w_i}(w) = 0$
- It is easy to see that

$$\nabla \hat{e}(w) = \frac{1}{N} X^t X w - \frac{1}{N} X^t Y = \hat{R} w - \hat{b}$$

### Solving the Linear Equations

- The optimal  $\hat{w}^*$  must verify  $\nabla \hat{e}(\hat{w}) = \hat{R} \hat{w} - \hat{b} = 0$ , where

$$\hat{R} = \frac{1}{N} X^t X, \quad \hat{b} = \frac{1}{N} X^t Y$$

- Over the original, non-centered data matrix we have

$$\hat{R} = \frac{1}{N} (X - \bar{X})(X - \bar{X})^t;$$

i.e.,  $\hat{R}$  is the **sample covariance matrix**

- If  $\hat{R}$  is **invertible**, we just solve the linear system  $\hat{R} \hat{w} = \hat{b}$

- And obtain the sample-dependent optimal  $\hat{w}^*$  as

$$\hat{w}^* = \hat{R}^{-1}\hat{b} = (X^t X)^{-1} X^t Y = \text{covar}(X)^{-1} \text{covar}(X, Y)$$

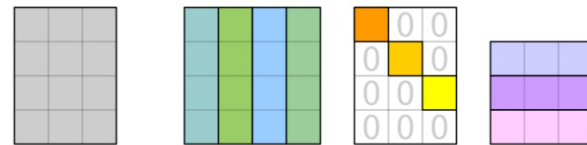
- $\hat{R}$  is oftent inverted through the **Singular Value Decomposition** (SVD) of the data matrix  $X$

## SVD

- The **Singular Value Decomposition** (SVD) of a  $N \times d$  matrix  $X$  is  $X = UDV^t$  where
  - $U, V$  are orthonormal matrices with dimensions  $N \times d$  and  $d \times d$
  - $D$  is a  $d \times d$  diagonal matrix with diagonal elements  $\delta_1 \geq \delta_2 \geq \dots$
- The columns of  $U$  and  $V$  are called the left and right singular vectors; the elements  $\delta_i$  are the singular values
- Then  $X^t X = VDU^t UDV^t = VD^2V^t$  and  $(X^t X)^{-1} = VD^{-2}V^t$ 
  - $D^{-2}$  is a diagonal matrix with  $(D^{-2})_{ii} = 1/D_{ii}^2$
- The columns of  $U$  and  $V$  are unique up to a sign change
  - So different packages may give seemingly different  $U$  and  $V$

## SVD II

- Sometimes  $U$  is given as an  $N \times N$  orthogonal matrix and  $D$  as an  $N \times d$  matrix with zeros in its last  $N - d$  rows
- We can represent graphically such a SVD decomposition as follows:



$$\begin{matrix} \mathbf{M} & = & \mathbf{U} & \mathbf{\Sigma} & \mathbf{V}^* \\ m \times n & & m \times m & m \times n & n \times n \end{matrix}$$

Taken from Wikipedia

- We may see this as  $M$  performing first a rotation with  $V^*$  (our  $V^t$ ), then feature dilations with  $\Sigma$  (our  $D$ ) and, finally, another rotation with  $U$ 
  - Actually both are a roto-reflections

## Gradient Descent MSE Minimization

- Computing the covariance matrix has a  $O(N \times d^2)$  cost and invert it has a  $O(d^3)$  cost

- For big data problems it may not be possible to solve analytically the normal equation  $\nabla \hat{e}(w) = 0$
- The simplest numerical alternative is **gradient descent**:
  - Starting from some random  $w^0$ , we iteratively compute
 
$$w^{k+1} = w^k - \rho_k \nabla \hat{e}(w^k) = w^k - \frac{\rho}{n_B} \left( \hat{X}_B^t \hat{X}_B w^k - \hat{X}_B^t Y \right)$$
 over a **mini-batch**  $B$  with  $n_B$  samples
    - Component wise:  $w_i^{k+1} = w_i^k - \rho_k \frac{\partial \hat{e}}{\partial w_i}(w^k)$
    - $\rho_k$  is the **learning rate**
- If  $w^k \rightarrow w^*$ , then  $\nabla \hat{e}(w^*) = 0$ 
  - Since our problems have obviously minima, this should be enough

### By the Way ...

- We have just built our first neural network: the linear regression architecture can be seen as a multilayer perceptron (MLP) with no hidden layers
- Thinking of this as an MLP seems far-fetched, as we are missing the hidden layers and the much more complicated gradient computation
- But the layer architecture is there and we also have learning as iterative error minimization, i.e., **training**
- In fact, we will see that MLPs perform linear regression on the last hidden layer outputs
- And the process from the input to the last hidden layer produces an enhanced feature set upon which the regression error should be better

### Measuring Model Fit

- First option: **Root Square Error**  $RSE = \sqrt{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}$
- OK, but how to compare this? We must always have a **base model** to benchmark our results
- Simplest “model”: a constant  $w_0$ , which yields the mean  $\bar{y} = \frac{1}{N} \sum_1^N y^p$ , with square error

$$\frac{1}{N} \sum (y^p - \bar{y})^2 = \text{Var}(y)$$

- We can compare our model against this base model by computing

$$\frac{\sum (y^p - \hat{y}^p)^2}{\sum (y^p - \bar{y})^2} = \frac{RSE^2}{\text{Var}(y)}$$

- The widely used  $R^2$  coefficient is simply  $R^2 = 1 - \frac{RSE^2}{\text{Var}(y)}$

### Regularization

- Our regression solution  $\hat{w}^* = (X^t X)^{-1} X^t Y$  won't work if  $X^t X$  is not invertible
  - For instance, when some features are correlated
- We could fix this working instead with  $X^t X + \alpha I$  for some  $\alpha > 0$ 
  - It is then always positive definite and we can thus invert it
- To make this practical, note that  $\hat{w}^* = (X^t X + \alpha I)^{-1} X^t Y$  minimizes

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- This is the **Ridge Regression** problem
  - Our first example of **regularization**, a key technique in Machine Learning
  - **All ML models must be regularized in some way**
- Important issue: how to find the right choice for  $\alpha$ ?

### Takeaways on Linear Regression

1. We introduced **supervised** models
2. We have reviewed the essentials of the **linear regression model** (always the first thing to try)
3. We have considered model estimation as a problem on **error minimization**
4. We have seen how to build linear models **analytically and numerically**
5. We have seen how to **measure model fit**
6. We have introduced **regularization**

## 1.3 Cross Validation, Bias and Variance

### Sample Bias and Variance

- All built models are **sample dependent**
- With several **independent** samples  $S_1, \dots, S_M$ , it is natural to use as our best final model the average of their associated  $\hat{f}_{S_m}(x)$  models, i.e.,

$$\frac{1}{M} \sum_1^M \hat{f}_{S_m}(x) \simeq E_S[\hat{f}_S(x)] = \hat{f}_N(x)$$

- Here we take the (theoretical) expectation  $E_S[\hat{f}_S(x)]$  over all samples  $S$  of size  $N$
- $\hat{f}_N(x) = E_S[\hat{f}_S(x)]$  is our ideal **best model**
- The **variance** of the  $\hat{f}_S(x)$  estimates is then

$$V_N(x) = E_S [(\hat{f}_S(x) - \hat{f}_N(x))^2]$$

### Bias Versus Variance

- Recall that our goal is to estimate  $\phi$  in our regression model  $y = \phi(x) + n$
- Ideally we would like to have a model such that


$$\hat{f}_N(x) - \phi(x) \simeq 0,$$

i.e., a model with small **bias**

- This should be achievable if we have rich, highly flexible models
- Or with essentially no regularization (or both!)
- But we would also like to have a robust model building procedure which results in a small **variance**  $V_N(x) \simeq 0$ 
  - This should be achievable if models are simple with few parameters
  - Or with more severe regularization if not
- Obviously both goals are contradictory to a large extent

### The Bias–Variance Tradeoff

- There is thus a **tradeoff** between bias (low for complex models) and variance (low for simple models)



Taken from *Hastie et al.*, p. 38

### Bias vs Variance in $k$ -NN Regression

- Recall that in  $k$ -NN regression we choose the  $k$  nearest neighbors  $x^{p_1}, \dots, x^{p_k}$  of  $x$  and estimate  $\hat{y} = \hat{y}(x)$  as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- Here the parameter that controls the tradeoff is just  $k$
- If  $k = N$ , the sample size, the  $N$ -NN estimator is just the mean:  $Y_N^{NN} = \bar{y}$ , with very small variance but large bias (it's an obviously bad model!!)
- If  $k = 1$ , the 1-NN estimator will have smaller bias but a large variance: changing the sample is very likely to change the sample point nearest to  $x$

### Bias vs Variance in Ridge Regression

- Recall that the Ridge Regression loss is

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- Here the parameter that controls the tradeoff is  $\alpha$
- If  $\alpha \gg 1$ , any non zero  $w$  implies a large regularization penalty
  - It is thus likely that  $w \simeq 0$  and the Ridge model reduces again to the mean  $\bar{y}$ , with small variance but possibly a large bias
- But if  $\alpha \simeq 0$ ,  $w$  can wander on the entire  $\mathbf{R}^d$ 
  - The bias will be then smaller, but the weights  $w_S$  and  $w_{S'}$  from different samples may be very different, particularly if the covariance is nearly singular

### Evaluating Expected Performance

- It is obvious that before we start applying a model, we should have a reasonably accurate idea of its performance in practice
- I.e., we want to estimate the model's **generalization performance**
- Estimating the generalization performance **only over the sample  $S$  used for training results in misleading error values**
- The preceding suggests to have  $M$  independent subsamples  $S_m$  and then
  - To compute  $\hat{f}_M(x) = \frac{1}{M} \sum_m \hat{f}_{S_m}(x) \simeq \hat{f}_N(x)$
  - To get the error estimate  $\hat{e} = \frac{1}{N} \sum_p (y^p - \hat{f}_M(x^p))^2$  over a new, **unseen** sample  $S' = \{(x^p, y^p)\}$
- But since usually we only have a single  $S$ , we “simulate” this by **Cross Validation** (CV) to get our first realistic generalization error estimates

### Cross Validation

- In Cross Validation (CV) we
  - Randomly split the sample  $S$  in  $M$  subsets  $S_1, \dots, S_M$
  - Work with  $M$  **folds**: pairs  $(S_m, S_m^c)$ , with
 
$$S_m^c = S - S_m = \cup_{i \neq m} S_i$$
  - Build  $M$  different models **using the  $S_m^c$  as training subsets**
  - Compute their errors  $e_m$  on the folds' **validation subsets  $S_m$**
  - Use these errors' average as a first estimate of the true model performance
- CV can and **must be used** in any model building procedure
  - Most data science packages have tools to simplify this
- We will also use CV to find **optimal model hyper-parameters** such as  $\alpha$  in Ridge Regression

### Grid Hyper-parameter Selection

- Consider for Ridge regression a hyperparameter range  $[\epsilon, A]$ , with  $0 < \epsilon < A$ 
  - In general  $\epsilon \simeq 0$ ,  $A \gg 1$
  - $\alpha = \epsilon$ : essentially no penalty, small bias and high variance
  - $\alpha = A$ : large penalty, small variance but high bias
- Select an  $L + 1$  point **grid**  $\{\alpha_0 = \epsilon, \alpha_1, \dots, \alpha_L = A\}$ 
  - The  $\alpha_j$  can be equi-spaced, log equi-spaced, random, ...
- At each  $\alpha_\ell$ 
  - Train  $M$  models on the  $S_m^c$  using the hyperparameter  $\alpha_\ell$
  - Average their test errors  $e_m$  on the  $S_m$  to get the error  $e(\alpha_\ell)$  at  $\alpha_\ell$
- And choose the final (hopefully) optimal hyperparameter  $\alpha^*$  as
 
$$\alpha^* = \arg \min_{0 \leq \ell \leq L} e(\alpha_\ell)$$
- $\alpha^*$  gives the model with the **best expected generalization among all possible  $\alpha$  choices**

### Takeaways on Bias, Variance and CV

1. We have stressed that **any model estimation is sample-dependent** and that this has to be controlled
2. We have introduced the **bias** and **variance** as the two key components of any model error
3. We have discussed **bias-variance trade-off**
4. We have introduced **Cross Validation** here as a tool to estimate a **model's generalization performance**
5. We have also introduced **Cross Validation** as a tool to estimate a **model's hyper-parameters**



## 1.4 Basic Classification

### Regression vs Classification

- Recall that in regression we have numerical continuous targets  $y$  and want our predictions  $\hat{y}$  to be as close to  $y$  as possible
- But in classification we have a finite number of labelled targets for which “selection by closeness” doesn’t make sense
- Natural alternative: select the **most probable** label given the pattern  $x$  we have just received
  - The concrete labels used for targets do not matter much anymore
  - Model learning should thus be “target” agnostic
  - And good probability estimates should be quite useful
- Let’s analyze this in an example

### A First Problem: Pima Indian Diabetes

- We want to diagnose whether a person may have diabetes from some clinical measures
- Features  $x$ : clinical measures
  - numPregnant
  - bloodPress
  - massIndex
  - age ...
- Target  $y$ : 0 (no diabetes), 1 (diabetes)
- Clear goal but perhaps too radical
- Better: try to estimate the probability  $P(1|x)$  of having diabetes depending on the features  $x$  we measure

### Classification Setup

- We have random patterns  $\omega$  from  $M$  classes,  $C_1, \dots, C_M$
- Over each pattern we “measure”  $d$  features  $x = x(\omega) \in \mathbb{R}^d$ 
  - $x$  inherits the randomness in  $\omega$  and becomes a random variable
- A  $\omega$  has a **prior probability**  $\pi_m$  of belonging to  $C_m$
- Inside each class  $C_m$  there is a **conditional class density**  $f(x|m)$  that “controls” the appearance of a given  $x$
- The  $\pi_m$  and  $f(x|m)$  determine the **posterior probability**  $P(m|x)$  that  $x$  comes from class  $C_m$

- **Intuition:** we should assign  $x$  to the class with the largest  $P(m|x)$ , that is, work with the classifier

$$\delta(x) = \arg \max_m P(m|x)$$

### Computing Posterior Probabilities I

- **Bayes rule:**  $P(B|A) = \frac{P(A \cap B)}{P(A)}$
- This requires to work with probabilities, not densities, but  $P(\{x\}) = P(m \cap \{x\}) = 0$  and

$$P(m|x) = \frac{P(m \cap \{x\})}{P(\{x\})} = \frac{0}{0} = \dots???$$

- But we can use the approximation

$$\begin{aligned} P(m|x) &\simeq P(m|B_r(x)) = \frac{P(C_m \cap B_r(x))}{P(B_r(x))} = \frac{P(B_r(x)|m)P(C_m)}{P(B_r(x))} \\ &= \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} = \pi_m \frac{\int_{B_r(x)} f(y|m) dy}{\int_{B_r(x)} f(z) dz} \end{aligned}$$

where we assume that features  $x$  are measured independently from classes  $m$

### Computing Posterior Probabilities II

- Remember the Fundamental Theorem of Calculus:  
if  $F(x) = \int_a^x f(y) dy$ ,

$$\lim_{\epsilon \rightarrow 0} \frac{1}{2\epsilon} \int_{x_0-\epsilon}^{x_0+\epsilon} f(y) dy = \frac{dF}{dx}(x_0) = f(x_0)$$

- In  $d$  dimensions it becomes

$$g(w) = \lim_{r \rightarrow 0} \frac{1}{|B_r(w)|} \int_{B_r(w)} g(z) dz$$

- Putting everything together, we arrive

$$\begin{aligned} P(m|x) &= \lim_{r \rightarrow 0} P(m|B_r(x)) = \pi_m \lim_{r \rightarrow 0} \frac{\int_{B_r(x)} f(y|m) dy}{\int_{B_r(x)} f(z) dz} \\ &= \pi_m \lim_{r \rightarrow 0} \frac{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(y|m) dy}{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(z) dz} = \frac{\pi_m f(x|m)}{f(x)} \end{aligned}$$

### The Obviously Optimal Classifier

- This suggests that a natural choice is the **classifier**  $\delta_B$

$$\begin{aligned} \delta_B(x) &= \arg \max_m P(m|x) = \arg \max_m \frac{\pi_m f(x|m)}{f(x)} \\ &= \arg \max_m \pi_m f(x|m) \end{aligned}$$

- With some extra work we can show that this **Bayes Classifier**  $\delta_B$  defines an optimal solution (in some precise sense) of the classification problem
- But ... this doesn't look too practical, for we do not know either  $\pi_m$  or (much harder)  $f(x|m)$

### Approximating the Bayes Classifier

- To define  $\delta_B$  we need to know the prior probabilities  $\pi_m$  and the prior densities  $f(x|m)$
- A reasonable choice for  $\pi_m$  is  $\hat{\pi}_m = \frac{N_m}{N}$ , where  $N_m$  is the number of patterns of  $C_m$  in the sample
- But effective multidimensional density estimates are rather difficult, because of the **curse of dimensionality**
  - Densities generalize histograms
  - Good histograms need accurate counts of elements nearby
  - But in high dimensions there won't be nearby elements!!
- Options:
  - Restrict possible density models: **logistic regression**
  - Assume no model and apply a **Nearest Neighbor (NN) strategy**

### The $k$ -NN Classifier

- Very simple: at any  $x$  consider the subset  $N_k(x)$  of its  $k$  closest sample points and
  - Let  $n_m(x)$  the number of elements of class  $m$  in  $N_k(x)$
  - Notice that  $0 \leq n_m(x) \leq k$
  - Define  $\delta_{kNN}(x) = \arg \max_m n_m(x)$
- That is,  $\delta_{kNN}(x)$  assigns  $x$  to the class that has more patterns in  $N_k(x)$
- We can partially justify this definition from a Bayesian point of view
- Assume that  $B_r(x)$  is the smallest ball that contains  $N_k(x)$  and consider the approximations
  - $P(B_r(x)|m) \simeq \frac{n_m(x)}{N_m}$
  - Similarly,  $P(B_r(x)) \simeq \frac{k}{N}$
  - And  $\pi_m \simeq \frac{N_m}{N}$

### $k$ -NN and the Bayes Classifier

- We then have

$$\begin{aligned}
 P(m|x) &\simeq P(m|B_r(x)) = \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} \\
 &\simeq \frac{N_m}{N} \frac{n_m(x)}{N_m} \frac{1}{\frac{k}{N}} = \frac{n_m(x)}{k}
 \end{aligned}$$

- Therefore  $\delta_{kNN}$  might be close to  $\delta_B$ , for

$$\begin{aligned}\delta_{kNN}(x) &= \arg \max_m n_m(x) = \arg \max_m \frac{n_m(x)}{k} \\ &\simeq \arg \max_m P(m|x) = \delta_B(x)\end{aligned}$$

### Some $k$ -NN Issues

- **Q1: How do we choose  $k$ ?** Using CV, of course
- There are no closed form solution and we have to deal again with the bias–variance tradeoff
  - Small variance with large  $k$ : if  $k = N$ ,  $k$ -NN classification returns the majority class
  - Small bias with small  $k$ : if  $k = 1$  a point very close to  $x$  should be in the same class
  - But also large variance: the nearest point to  $x$  in another sample may well belong to a different class
- **Q2: Is  $k$ -NN always meaningful?**
  - Well, only if classes are concentrated and far from each other
  - In other words, when the classification problem is easy!!

### Measuring Classifier Accuracy

- Consider a two class problem with labels  $y = 0, 1$
- We will call patterns with label 1 **positive** and those with label 0 **negative**
  - Usually the positive patterns are the interesting ones: sick people, defaulted loans, . . .
- Let  $\hat{y} = \hat{y}(x)$  the label predicted at  $x$ ; we say that  $x$  is a
  - **True Positive (TP)** if  $y = \hat{y} = 1$
  - **True Negative (TN)** if  $y = \hat{y} = 0$
  - **False Positive (FP)** if  $y = 0$  but  $\hat{y} = 1$
  - **False Negative (FN)** if  $y = 1$  but  $\hat{y} = 0$
- These values are combined in many performance measures

### Classifier Metrics

- The classifier **accuracy** is  $acc = \frac{TP+TN}{N}$ 
  - The first thing to measure but ...
- It may not be too significant when the number  $N_0$  of negatives is  $\gg N_1$ , the number of positives; then
  - The classifier  $\delta(x) = 0$  will have a high accuracy  $N_0/N \simeq 1$

- But it will also be useless!!
- First variant: Precision, Recall
  - **Recall:**  $TP/(TP + FN)$ , i.e., the fraction of positives detected
  - **Precision:**  $TP/(TP + FP)$ , i.e., the fraction of true alarms issued
- Recall measures how many positive cases we recover, i.e., how effective is our method
- Precision measures the effort we need for that, i.e., its efficiency
- Ideally: high recall, high precision (i.e., effective and efficient!!)
- Another way of presenting the performance is through the **confusion matrix**

### The Confusion Matrix

- Standard layout

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

- Other layouts:
  - **Positives (with label 1) at bottom** (as done in `confusion_matrix` of `sklearn`)
  - Predicted values in rows, real values in columns
- Accuracy and the confusion matrix can be extended to multi-class problems
  - But not other concepts such as recall or precision

### Linear Regression for Classification?

- $k$ -NN Classifier is simple but also crude; have to look elsewhere
- Building a regression model with targets given by some coding of class labels usually doesn't make sense
- However, for a binary 0–1 response, it can be shown that the  $w_0 + w \cdot x$  obtained using linear regression is in fact an estimate of  $P(1|x)$ 
  - We may thus fix a threshold  $\delta_0$  and decide 0 if  $w_0 + w \cdot x < \delta_0$  and 1 otherwise
  - However, we may end up with probability estimates less than 0 or bigger than 1!!!

- Better idea: try to transform the linear output  $w_0 + w \cdot x \in (-\infty, \infty)$  into a probability  $P(1, x) \in (0, 1)$ 
  - Logistic regression does this through the **logistic** or **sigmoid** function

### Logistic Regression (LogR)

- We assume

$$P(1|x) = P(1|x; w_0, w) = \frac{1}{1 + e^{-(w_0 + w \cdot x)}}$$

- Then  $0 \leq P(1|x) \leq 1$  for any  $x$

- We then have

$$P(0|x) = 1 - P(1|x) = \frac{e^{-(w_0 + w \cdot x)}}{1 + e^{-(w_0 + w \cdot x)}} = \frac{1}{1 + e^{w_0 + w \cdot x}}$$

- Notice that if  $w_0 + w \cdot x = 0$ ,  $P(1|x) = P(0|x) = 0.5$

- The ratio  $\frac{P(1|x)}{P(0|x)} = e^{w_0 + w \cdot x}$  is called the **odds** of  $x$  and its log the **log odds** or **logit**
- Thus, the basic assumption in LogR is that the **logit is a linear function**  $w_0 + w \cdot x$  of  $x$
- We have the model  $f(x; w)$ ; we need a **loss** function  $L(w_0, w)$  to minimize for which we use the sample's **likelihood**

### Estimating $w_0^*, w^*$

- Assume a single sample  $x, y$  and two possible model coefficients  $w_0, w$  and  $w'_0, w'$
- Denoting by  $p = P(y|x; w_0, w)$  and  $p' = P(y|x; w'_0, w')$ , it is clear that we should prefer  $w_0, w$  if  $p > p'$  and  $w'_0, w'$  if not
  - In other words, we prefer the coefficients that give a **higher posterior probability**
- For an **independent** sample  $S = \{(x^p, y^p)\}$ , its joint probability under a posterior model  $p = P(y|x, w_0, w)$  is

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And, again, given two possible model coefficients  $w_0, w$  and  $w'_0, w'$ , **we should prefer  $w_0, w$  iff**

$$P(Y|X; w_0, w) > P(Y|X; w'_0, w')$$

### Sample's Likelihood

- Therefore, we can estimate the optimal  $w_0^*, w^*$  as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- By the independence assumption we have

$$\begin{aligned}
 P(Y|X; w_0, w) &= \prod_{p=1}^N P(y^p|x^p; w_0, w) \\
 &= \left\{ \prod_{y^p=1} P(1|x^p) \right\} \left\{ \prod_{y^p=0} P(0|x^p) \right\} \\
 &= \prod_{p=1}^N P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}
 \end{aligned}$$

because

- If  $y^p = 1$ ,  $P(1|x^p) = P(1|x^p)^{y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$ , and
- If  $y^p = 0$ ,  $P(0|x^p) = P(0|x^p)^{1-y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$

### Max Log-Likelihood Estimation

- The **log-likelihood** of  $w_0, w$  given  $S$  is then

$$\begin{aligned}
 \ell(w_0, w; S) &= \log P(Y|X; w_0, w) \\
 &= \sum_p \{y^p \log p(1|x^p) + (1 - y^p) \log p(0|x^p)\} \\
 &= \sum_p y^p \log \frac{p(1|x^p)}{p(0|x^p)} + \sum_p \log p(0|x^p) \\
 &= \sum_p y^p (w_0 + w \cdot x^p) - \sum_p \log(1 + e^{w_0 + w \cdot x^p})
 \end{aligned}$$

- We can thus estimate the optimal  $\hat{w}_0^*, \hat{w}^*$  as

$$\hat{w}_0^*, \hat{w}^* = \arg \min_{w_0, w} -\ell(w_0, w; S)$$

- Extra bonus:  $-\ell$  is a convex differentiable function of  $(w_0, w)$  and, thus, it is enough to solve  $\nabla \ell(w_0, w) = 0$

### Newton-Raphson Solution

- However,  $\nabla \ell(w_0, w) = 0$  doesn't admit a closed form solution, but only an iterative, numerical one
- We apply the **Newton-Raphson** iterative method, here equivalent to the general Newton method for function minimization
- For a general weight vector  $w$ , and starting with an initial random  $w^0$ , Newton's iterations are

$$w^{k+1} = w^k + (\mathcal{H}_\ell(w^k))^{-1} \nabla \ell(w^k)$$

- $\mathcal{H}_\ell(w^k)$  denotes the **Hessian** of  $\ell$  at  $w^k$ , which may or may not be invertible

- Everything is fine if the  $w^k$  are close enough to the optimum  $w^*$  but far away things may get tricky
- Just as before, we can add a regularization term  $\frac{\alpha}{2} \|w\|^2$  to avoid invertibility problems
- The iterations in Logistic Regression are again typical of many of the model building methods used in Machine Learning

### Recap: Learning in ML

- The general approach to **learning** is the following:
  - A **model**  $f(x; w)$  is chosen
  - Given a sample  $S = \{(x^1, y^1), \dots, (x^N, y^N)\}$ , we define a **sample dependent loss function**

$$L(w) = L(w|S) = L(y^1, \dots, y^N, f(x^1; w), \dots, f(x^N; w))$$

- $L(w)$  is often minimized from some  $w^0$  by **iterations**

$$w^{k+1} = w^k - \rho_k G(w^k, S)$$

with  $\rho_k$  a **learning rate** and  $G$  some vectorial function

- When  $G(w) = \nabla L(w)$  we have **gradient descent**
- When  $G(w) = \mathcal{H}(w)^{-1} \nabla L(w)$  we obtain **Newton's method**
- When the entire sample  $S$  is used at each iteration, we speak of **batch learning**
- When only single patterns  $(x^p, y^p)$  or small subsamples are used, we speak of **on-line** or **mini-batch learning**

### Takeaways on Classification

1. We have introduced the classification problem as one of **computing posterior probabilities**
2. We have found the **optimal Bayes classifier** and **approximated it by  $k$ -NN**
3. We have introduced several measures of **classifier performance**
4. We have introduced **Logistic Regression** and the numerical minimization of its (minus) log-likelihood
5. We have put Logistic Regression modeling in the general framework of Machine Learning models


## 2 Multilayer Perceptrons

### 2.1 Rosenblatt's Perceptrons

#### Basic Neural Models

- Basic model: Ramón y Cajal's neuron (1900)






- Basic behavior: the neuron either fires or stays at rest depending basically on its inputs
- This translates into an electrical nerve impulse or **action potential**
- The brain has about  $10^{11}$  neurons
  - Each one has about 7,000 connections
  - These connections are often **recurrent**

### Hodgkin–Huxley

- They developed (circa 1935) the first electro-physiological model to describe the generation and propagation of action potentials in neurons
- They used the giant squid axon for this




From Wikipedia's [Squid Giant Synapse](#)

### Electronic Neuron

- Electronic version: McCulloch–Pitts (1940)
  - Since  $I = \frac{V}{R} = wV$ , with  $w$  the conductance, the McCulloch–Pitts neuron output is

$$H \left( \sum_{j=1}^d w_j V_j + I \right) = V$$

with the Heaviside function  $H$  ensuring a 0–1 output



- The **perceptron** is a mathematical formulation of the McC-P neuron where **the conductances are to be learned**

### Rosenblatt's Perceptron

- Given a sample  $S = \{(x^p, y^p)\}$  with  $y^p = \pm 1$ , **Rosenblatt's Perceptrons** (PCPs) are linear machines  $w \cdot x$  such that

$$\begin{aligned} w \cdot x^p &> 0 & \text{if } y^p &= 1 \\ w \cdot x^p &< 0 & \text{if } y^p &= -1 \end{aligned}$$

- This can be seen as a **homogeneous** classification problem
- In more compact form we want for all  $p$

$$y^p w \cdot x^p > 0$$

- If such a **separating**  $w$  exists, it can be computed in many ways
- But Rosenblatt's goal was to **learn**  $w$ , i.e., to arrive at a separating  $w$  by repeatedly examining the  $(x^p, y^p)$  and adjusting  $w$  if necessary

### Rosenblatt's Delta Rule

- The **Delta Rule** algorithm

```

Start with  $w = 0$ 
While stopping condition not met:
  Get a new pattern  $(x^p, y^p) = (x^{p(t)}, y^{p(t)})$ 
  if  $y^p w \cdot x^p \leq 0$  then:
     $w = w + y^p x^p$ 

```

- Writing  $w$  for the new vector and  $w'$  for the previous one, notice that

$$y^p w \cdot x^p = y^p w' \cdot x^p + \|x^p\|^2 > y^p w' \cdot x^p$$

### Novikov's Theorem

- It states that if the sample  $S$  is **linearly separable**, the Delta rule yields a separating  $w$  in a **finite number of steps**
  - I.e., it **learns!!**
- More precisely, if for some separating hyperplane  $w$  with  $\|w\| = 1$  we have  $\min_p y^p w \cdot x^p = \gamma > 0$ , then **a separating  $w'$  can be found after at most  $T = R^2/\gamma^2$  wrong classifications, where  $R = \max_p \|x^p\|$**
- Notice that  $\gamma$  is the **margin** of the  $w$  hyperplane. Thus, if


$$\gamma^* = \max_{\|w\|=1} m(w, S) = \max_{\|w\|=1} \min_p y^p w \cdot x^p$$

is the **maximum margin**, we will need at most  $R^2/(\gamma^*)^2$  iterations to train a Perceptron

- The maximum margin hyperplane gives the best (smallest) Novikov bound
- And problems with smaller margins require more iterations


### Margins and Generalization

- Intuitively, when the margin is large, the problem is easier



### Distance to a Line

- The margin of  $w$  is just the smallest distance between sample points and the hyperplane defined by  $w$
- This extends to the multidimensional case: if  $\|w\| = 1$  and  $b = 0$ , the distance of  $x$  to the hyperplane defined by  $w$  is  $|w \cdot x| = yw \cdot x$



- Thus the Delta Rule **corrects weight vectors with negative margins**

#### Novikov's Proof Sketch

- Assume  $x^t$  is the  $t$ -th pattern such that  $y^t w^{t-1} \cdot x^t \leq 0$
- Then  $\|w^t\|^2 = \|y^t x^t + w^{t-1}\|^2 \leq \|x^t\|^2 + \|w^{t-1}\|^2$  and, thus,

$$\|w^t\|^2 \leq \|x^t\|^2 + \dots + \|x^1\|^2 + \|w^0\|^2 \leq tR^2$$

if  $w^0 = 0$  and  $\|x^p\| \leq R$  for all  $p$

- Also, if  $w'$  is a 1-norm separating hyperplane i.e.,  $\min_p y^p w' \cdot x^p = \gamma > 0$  for all  $p$ , we have

$$w^t \cdot w' = (y^t x^t + w^{t-1}) \cdot w' \geq \gamma + w^{t-1} \cdot w' \geq \dots \geq t\gamma + w^0 \cdot w'$$

- Putting both together we have for all  $t$  for which  $x^t$  is erroneous

$$t\gamma \leq w^t \cdot w' \leq \|w^t\| \leq \sqrt{t}R$$


which implies  $t \leq R^2/\gamma^2$

- Thus, the Delta Rule will stop after at most  $R^2/\gamma^2$  iterations

#### Problems with Perceptrons

- Rosenblatt's Perceptron caused a big stir on the feasibility of **Conexionist Artificial Intelligence** but soon its first winter came
- Observation 1: A PCP can learn (obviously) only linearly separable problems; for instance, it can learn the AND and OR predicates
- But a **PCP cannot learn the XOR predicate**

#### Dichotomies



- Observation 2: for PCPs to be useful, **linearly separable problems must be frequent**
- Q: how frequent are they?
- A **dichotomy** for  $S = \{x^1, x^2, \dots, x^N\}$  is any separation of  $S$  in two classes
- It corresponds to any assignment of the  $\pm 1$  values to the  $x^p$
- The total number of dichotomies is  $2^N$
- Thus, the **total number of 2-class problems over  $S$  is  $2^N$**

### Cover's Theorem


- We have to compare the total number of dichotomies with the number of dichotomies that are **linearly separable**
- The points in a sample  $S$  are said to be in **general position** if there are not  $d + 1$  patterns in a  $d - 1$ -dimensional hyperplane
- **Theorem** If  $S$  is in general position, the number  $L(N, d)$  of linearly separable dichotomies is

$$L(N, d) = \begin{cases} 2^N & \text{si } N \leq d + 1 \\ 2 \sum_{i=0}^d \binom{N-1}{i} & \text{si } N \geq d + 1 \end{cases}$$


### Points in General Position

- Consider  $d = 2$ ,  $3 = d + 1$  points and a  $1 = d - 1$ -dimensional hyperplane

### Points Not in General Position



- Consider now  $d = 2$  and  $3 = d + 1$  points **not** on a  $1 = d - 1$ -dimensional hyperplane (i.e., a line)



Thus ...

- Notice that for  $d$  fixed,

$$\frac{L(N, d)}{2^N} \rightarrow 0$$

as  $N \rightarrow \infty$

- In practice we can expect  $N \gg d$  and the fraction of separable dichotomies will be very small
- Thus, linear PCPs will not be useful

- But a **one hidden layer** PCP can solve XOR and, in fact, **separate any convex region** from the rest of the space
- And a **two hidden layer** PCP can separate any **polyhedral region** and hence **solve any classification problem**
- The same is essentially true for regression problems
- But: **no algorithm is known to learn these 2-hidden layer PCPs**

### Why Not Try to Minimize Some Loss?

- Rosenblatt's Perceptrons are natural **online** machines:
  - Patterns appear sequentially and each may change the current weights
- We can also define a perceptron loss over a given sample as

$$L(w|S) = - \sum_{\{p: y^p w \cdot x^p < 0\}} y^p w \cdot x^p$$


- The local error would thus be  $-y^p w \cdot x^p$  when  $y^p w \cdot x^p < 0$  and 0 otherwise
- Since  $\nabla_w (-y^p w \cdot x^p) = -y^p x^p$ , the Delta Rule can be seen as a kind of local gradient descent on the loss  $L(w|S)$ :

$$w^t = w^{t-1} + y^{p(t)} x^{p(t)} = w^{t-1} - \nabla_w \left( -y^{p(t)} x^{p(t)} \cdot w \right)$$

- However  $L(w|S)$  is not differentiable and it is not clear what to make of  $\nabla_w L(w|S) \dots$

### What Can We Do?

- First option: work with linear models but somehow ensure that  $d \gg N$ 
  - Transform the original  $d$ -dimensional features  $x$  in new  $D$ -dimensional ones such that  $D \gg d$
  - We overcome then that, in practice, we usually have  $N \gg d$
  - This is the goal of kernel-based Support Vector Machines
- Second option: build trainable PCPs replacing the Heaviside function
  - Work with differentiable PCP transfer function
  - Transform PCP training into a differentiable optimization problem
  - This is what we will do next
- Before doing so, note that **Rosenblatt's PCPs can be seen as the first Machine Learning model**
- And modern ML started around 1985 along the PCP's ideas




## 2.2 Classical MLPs

### Basic Neural Models

- Basic model: Ramón y Cajal's neuron (1900)
- Basic behavior: the neuron either fires or stays at rest depending basically on its inputs
- The brain has about  $10^{11}$  neurons with each one having about 7,000 connections, often recurrent

### Hodgkin–Huxley

- They developed (circa 1935) the first model to describe the generation and propagation of electrical **action potentials** in neurons
- They used the giant squid axon for this




From Wikipedia's [Squid Giant Synapse](#)

### McCulloch–Pitts



- Idealized electronic version of a neuron's working (1943)



- Taking weights  $w = 1/R$  as conductances,  $x$  as potentials and  $wx = \frac{x}{R}$  as intensities, the McCulloch–Pitts neuron outputs a potential  $x_j$

$$x_j = H \left( \sum_{k=1}^N w_{jk} x_k + I_j \right)$$


- The Heaviside function  $H$  ensures a 0, 1 output
- I.e., the neuron fires or does not fire

### Basic Questions

- Q1: How to adjust the  $w_j$  and  $I$  values? I.e., how to “learn” them?
  - It will depend on the problem at hand but it is relatively easy for isolated neurons
- Q2: How to model and reproduce the joint behavior of groups of neurons?
  - Related to the previous questions but quite difficult!!
    - \* It is very hard to measure the joint behavior of groups of neurons
    - \* It is very hard to take into account the recurrence present in real neurons
- In Artificial Neural Networks (and in ML) one considers just Q1 and (usually) outside any neuro-computational framework

### The Perceptron

- Proposed by **Frank Rosenblatt** (1957) as a first step to build an artificial retina



- Key step: separate two classes of vectors  $x$  labelled with a target  $y = \pm 1$
- Want a weight vector  $w$  such that  $w \cdot x > 0$  if  $y = 1$  and  $w \cdot x < -1$  if  $y = -1$
- In a unified form we want  $y w \cdot x > 0$
- First success: this can be done by **learning**

### The Delta Rule

- Connectionist paradigm: learning means that weights change as new inputs are received
  - They go through a series of steps  $w = w^k$  as they receive the  $k$ -th pair  $x^p, y^p$
  - In the Perceptron's case, change  $w$  when  $y^p w \cdot x^p \leq 0$

- Perceptron's adaptive steps:

```

w = 0
while there remain wrongly classified patterns:
    get a new pattern (x_p, y_p)
    if y_p * w.dot(x_p) <= 0:
        # w's response was wrong so we change it
        w = w + y_p * x_p

```


- Then the Perceptron **learns**: if the patterns can be separated, after a finite number of steps we arrive at a  $w$  **that separates the patterns**
- But ...

### The Connectionist Winter

- Perceptron learning caused a lot of excitement at the time
- However, it could not solve simple separation problems such as the XOR:

01 <sub>1</sub>	11 <sub>0</sub>
00 <sub>0</sub>	10 <sub>1</sub>

- This can be solved adding more layers to the basic Perceptron




i.e., to work with **Multilayer Perceptrons (MLPs)**

- But **no algorithm to make them learn** was found, and interest in PCPs decayed

### MLP's Rebirth

- New proposal (circa 1985): keep the architecture but look to MLP learning as a form of **error minimization**
- MLP's General layout:
  - An input layer (input)
  - One or several hidden layers
  - One output layer
- Feedforward connections only



- Example: [TensorFlow Playground](#)

### MLP Connections

- **No feedback or lateral connections**
- **Fully connected layers**
- **Linear weight connections** between units and (usually) **non linear activations** inside each unit
  - Combined effect of successive layers: potentially highly non-linear transformation
- General processing: layered and feedforward
- In practice (1990s), one hidden layer and only sometimes two
- Later (around 2010): Deep Networks with “many” (from 3 to 10) layers
- Now: almost anything!

### Unit Activation and Output

- The **activations**  $a_i^h$  of a unit in layer  $h$  receives the **outputs**  $o_j^{h-1}$  from processing in the previous layer

$$a_i^h = \sum_{j=1}^{n_{h-1}} w_{ij}^h o_j^{h-1} + b_i^h,$$

- In matrix/vector form:

$$a^h = W^h o^{h-1} + b^h$$

- **Output**  $o_i^h$  of a unit: **non linear processing**  $o_i^h = \varphi(a_i^h)$  of its activation
- In vector form:

$$o^h = \varphi(a^h),$$


where  $\varphi$  is applied over each unit

### Activation Functions

- Choices for  $f$ :
  - Heaviside (in the very first Rosenblatt's Perceptrons):  $\varphi(a) = 0$  if  $a \leq 0$ ,  $\varphi(a) = 1$  if  $a > 0$
  - **Identity/linear**:  $\varphi(a) = a$
  - **Sigmoid**:
 
$$\varphi(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$
  - **Hyperbolic tangent**:
 
$$\varphi(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$
  - **Rectified Linear Units (ReLU)**:  $\varphi(a) = r(a) = \max(0, a)$

### Sigmoid, Hyperbolic Tangent, ReLUs

- Sigmoid and tanh: smooth version of Heaviside step function
- For ReLUs  $r'(x)$  is either 0 or 1 (hoping  $x = 0$  never happens!!)
  - Many gradient elements will go to 0 and many units will have constant activations no matter their input



From [Stanford's UFLDL Tutorial](#)

- Usual activation choices:

- For hidden units: hyperbolic tangent and, currently, ReLU
- Linear outputs for regression, sigmoid/soft max outputs for classification

### The Simplest MLP I

- The Single Hidden Layer (SHL) MLP
  - $D$  inputs (determined by the problem at hand)
  - One hidden layer with  $H$  units (number to be chosen) and tanh (or what you prefer) activation
  - One linear or sigmoid output (according to the problem at hand)
- Input–hidden processing: denoting inputs by  $x$  and the hidden unit output as  $o^H$ ,

$$o_h^H = \tanh \left( b_h^H + \sum_{j=1}^D W_{hj}^H x_j \right)$$

- In matrix/vector form:  $o^H = \tanh (W^H x + b^H)$

### The Simplest MLP II

- Hidden–output processing: since targets are 1–dimensional, we have for the outputs  $\hat{y}$

$$\hat{y} = \sum_{h=0}^H w_h^O o_h^H + b^O,$$

- In vector form:  $\hat{y} = w^O \cdot o^H + b^O$
- That is, a **linear model** on the last hidden layer outputs  $o^H$
- Global process:

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + \sum_h w_h^O \tanh \left( b_h^H + \sum_j W_{hj}^H x_j \right)$$

- Or in matrix/vector form

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + w^O \cdot \tanh (b^H + W^H x)$$

- Or in Python `y_pred = np.tanh(x.dot(w_H.T) + b_H).dot(w_O.T) + b_O`

### MLPs and Universal Approximation

- Recall that the regression goal is to get  $f \simeq \phi$
- We say that  $\mathcal{F} = \{f(x; \mathcal{W})\}$  is a **Universal Approximation Family** over a domain  $\mathcal{R}$  if

For any  $\epsilon > 0$  and any reasonable  $\phi$ , we can find an  $f(x; \mathcal{W}_{\phi, \epsilon})$  s.t.

$$\int (\phi(x) - f(x; \mathcal{W}_{\phi, \epsilon}))^2 p(x) dx \leq \epsilon$$

- Notice that Universal Approximation is just what we need in regression
- In fact a **Single Hidden Layer MLP with enough hidden units is an effective universal approximator**
- But we have to be able to build them

### MLP Error Function

- MSE is the standard error function for regression MLPs

$$\begin{aligned} e(\mathcal{W}) &= \frac{1}{2} E_{x,y} [(y - f(x; \mathcal{W}))^2] = E_{x,y} [e^\ell(x, y; \mathcal{W})] \\ &= \int e^\ell(x, y; \mathcal{W}) p(x, y) dx dy \end{aligned}$$

with  $e^\ell(x, y; \mathcal{W})$  denotes the **local error**

$$e^\ell(x, y; \mathcal{W}) = \frac{1}{2} (y - \hat{y})^2 = \frac{1}{2} (y - f(x; \mathcal{W}))^2$$

### MSE Gradient

- The general idea would be to obtain  $\mathcal{W}^*$  as a solution of  $\nabla e(\mathcal{W}) = 0$ , where we have

$$\begin{aligned} \nabla e(\mathcal{W}) &= E_{x,y} [\nabla_{\mathcal{W}} e^\ell(x, y; \mathcal{W})] \\ &= E_{x,y} [\nabla_{\mathcal{W}} f(x; \mathcal{W}) (f(x; \mathcal{W}) - y)] \end{aligned}$$

for we have

$$\begin{aligned} \nabla_{\mathcal{W}} e^\ell(x, y; \mathcal{W}) &= -(y - f(x; \mathcal{W})) \nabla_{\mathcal{W}} f(x; \mathcal{W}) \\ &= \nabla_{\mathcal{W}} f(x; \mathcal{W}) (f(x; \mathcal{W}) - y) \end{aligned}$$

- We have therefore two tasks:
  - Compute  $\nabla e$
  - Exploit it to build MLPs
- We will exploit  $\nabla e(\mathcal{W})$  through **optimization methods** after we compute it


### SHL Forward Pass I

- We need thus to compute  $\nabla e$ , for which we need to compute quite a few auxiliary variables
- We apply the preceding to a single hidden layer (SHL) MLP with
  - A single output unit and input-to-hidden weight matrix  $W^H = (W_{hj}^H)$  and bias  $b^H$  vector and
  - A hidden-to-output weight vector  $w^O = (w_1^O, \dots, w_H^O)$  and scalar bias  $b^O$

- Recall that the forward pass can be computed as follows
  - $a^H = W^H x + b^H, o^H = \varphi(a^H),$
  - Or unit-wise:  $a_h^H = \sum W_{hi}^H x_i + b_i^H, o_h^H = \varphi(a_h^H)$
  - $y = w^O \cdot o^H + b^O = \sum_h w_h^O o_h^H + b^O$
- Straightforward to program

### SHL Forward Pass II

- Graphically we have the following scheme:




From [Sebastian Raschka's A Visual Explanation of the Back Propagation Algorithm for Neural Networks](#), KDnuggets

- We turn next to the gradient computation
  - Do we need to work on that? Current Deep NN backends compute the error gradient “automatically”
  - But, as Andrej Karpathy says, [Yes, you should understand backprop](#)

### Backprop's Basic Idea

- Key idea: use the activations  $a_i^h$  as intermediate variables



- This suggests to apply the chain rule with the  $a_h$  in the next layer as intermediate variables
- Let  $w_{ij}^h$  be the weight connecting unit  $j$  in layer  $h - 1$  with unit  $i$  in layer  $h$ ; then

$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1}$$

### Backprop's Details

- If we have already computed the generalized errors  $\delta_k^{h+1} = \frac{\partial e^\ell}{\partial a_k^{h+1}}$  at the  $h + 1$  layer, we have

$$\begin{aligned} \frac{\partial e^\ell}{\partial a_i^h} &= \sum_k \frac{\partial e^\ell}{\partial a_k^{h+1}} \frac{\partial a_k^{h+1}}{\partial a_i^h} = \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial a_i^h} \\ &= \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial o_i^h} \frac{\partial o_i^h}{\partial a_i^h} = \left[ \sum_k \delta_k^{h+1} w_{ki}^{h+1} \right] \sigma'(a_i^h) \end{aligned}$$

- It has a  $O(1)$  cost, as we can precompute the bracketed term

- Thus we start at the easy partials  $\frac{\partial e^\ell}{\partial a_i^L}$  and  $\frac{\partial e^\ell}{\partial w_{ij}^L}$  at the last layer  $\mathcal{L}$  and go backwards

### SHL Gradient Backprop I

- We refer to  $\frac{\partial e^\ell}{\partial a_i} = \delta_i$  as the **generalized error**
- In the output layer  $e^\ell = \frac{1}{2}(y - \hat{y})^2$  and  $a^O = \hat{y}$ , and thus,

$$\delta^O = \frac{\partial e^\ell}{\partial a^O} = \frac{\partial e^\ell}{\partial \hat{y}} = \hat{y} - y$$

i.e., the generalized error  $\delta^O$  is here the standard error

- Therefore, in the output layer we have

$$\frac{\partial e^\ell}{\partial w_h^O} = (\hat{y} - y) \frac{\partial a^O}{\partial w_h^O} = (\hat{y} - y) o_h^H$$

### SHL Gradient Backprop II

- In the hidden layer we **backpropagate** the error  $\delta^O = \hat{y} - y$ :

$$\delta_h^H = \frac{\partial e^\ell}{\partial a_h^H} = \frac{\partial e^\ell}{\partial a^O} \frac{\partial a^O}{\partial a_h^H} = \delta^O \frac{\partial a^O}{\partial a_h^H} = (\hat{y} - y) \frac{\partial a^O}{\partial a_h^H}$$

- And since  $y = a^O = \sum_h w_h^O o_h + b^O = \sum_h w_h^O \varphi(a_h^H) + b^O$ ,

$$\frac{\partial a^O}{\partial a_h^H} = \frac{\partial a^O}{\partial o_h^H} \frac{\partial o_h^H}{\partial a_h^H} = \frac{\partial a^O}{\partial o_h^H} \varphi'(a_h^H) = w_h^O \varphi'(a_h^H)$$


- Moreover,  $a_h^H = \sum_j w_{hj}^H x_j + b_h^H$ , and

$$\frac{\partial e^\ell}{\partial w_{hj}^H} = \frac{\partial e^\ell}{\partial a_h^H} \frac{\partial a_h^H}{\partial w_{hj}^H} = \delta_h^H x_j = (\hat{y} - y) w_h^O \varphi'(a_h^H) x_j$$

### SHL Gradient Backprop III



- Graphically we have the following scheme:



From [Sebastian Raschka's A Visual Explanation of the Back Propagation Algorithm for Neural Networks](#), KDnuggets

### Backprop for Other Layers

- Once we analyze the SHL Backprop, we see that it can be extended easily to other layers
- Let's  $w_{ij}^h$  be the weight connecting unit  $j$  in layer  $h - 1$  with unit  $i$  in layer  $h$ ; then

$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1}$$

- And if we have already computed the generalized errors  $\delta_k^{h+1} = \frac{\partial e^\ell}{\partial a_k^{h+1}}$ , we have

$$\begin{aligned} \frac{\partial e^\ell}{\partial a_i^h} &= \sum_k \frac{\partial e^\ell}{\partial a_k^{h+1}} \frac{\partial a_k^{h+1}}{\partial a_i^h} = \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial a_i^h} \\ &= \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial o_i^h} \frac{\partial o_i^h}{\partial a_i^h} = \left[ \sum_k \delta_k^{h+1} w_{ki}^{h+1} \right] \sigma'(a_i^h) \end{aligned}$$

with a  $O(1)$  cost, as we can precompute the bracketed term

### Neural Network Classification

- Recall that our Logistic Regression model gave us estimate of  $P(0|x; w_0, w)$ ,  $P(1|x; w_0, w)$
- We then wrote the likelihood of a given sample as

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And then derived the optimum parameters  $w_0^*, w^*$  as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- Notice that we can see the output of a logistic regression model as that of an **NN with no hidden layers and sigmoid outputs**

**Neural 2-class Classification**

- We can follow the same path for a 2 class classification problem changing that output for the one

$$f(x; \mathcal{W}) = w_0 + w \cdot \Phi(x, \mathcal{W}^H)$$

with  $w_0, w$  the weights from the last hidden layer (LHL) to the output and  $\Phi(x, \mathcal{W}^H)$  the map from inputs to the LHL outputs

- And now derive the optimum parameters  $\mathcal{W}^*$  as

$$\begin{aligned} \mathcal{W}^* &= \arg \min_{\mathcal{W}} - \sum_{p=1}^N \log P(y^p | x^p; \mathcal{W}) \\ &= \arg \min_{w_0, w, \mathcal{W}^H} - \sum_p y^p (w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)) + \\ &\quad \sum_p \log(1 + e^{w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)}) \end{aligned}$$

- Forward and backpropagation are also straightforward here
- We change things a little bit for  $K$  class problems

**MLPs for  $K$ -class Classification**

- We consider an input layer and a number of hidden layers
- Targets are the 1-hot encodings of class labels, so we use  $K$  outputs
- We want the MLP's  $k$ -th output to estimate the posterior  $P(k|x)$
- The natural output layer activation is thus the **softmax function**

$$\sigma_j(o) = F_j(x; \mathcal{W}) = \frac{e^{w_j \cdot o}}{\sum_1^K e^{w_k \cdot o}}$$

- Here  $o$  are the last hidden layer outputs and  $w_j$  is the weight vector in  $\mathcal{W}$  which connects the last hidden layer with the network's  $j$  output
- Moreover  $\sum_j F_j(x; \mathcal{W}) = 1$  and we assume  $P(j|x) \simeq F_j(x; \mathcal{W})$
- For two classes this becomes  $\sigma_1(o) = \frac{e^{w_1 \cdot o}}{e^{w_0 \cdot o} + e^{w_1 \cdot o}} = \frac{1}{1 + e^{(w_0 - w_1) \cdot o}}$ 
  - We thus get the **sigmoid activation** of Logistic Regression

**Multiclass Log-Loss NN Classification**

- We change slightly the preceding for a  $K$  class classification problem
- We define an MLP  $F(x; \mathcal{W}) \in \mathbf{R}^K$  with softmax outputs
  - Recall that  $0 \leq F_k(x; \mathcal{W}) \leq 1$  and  $\sum_k F_k(x; \mathcal{W}) = 1$

- And also we assume that  $P(k|x) \simeq F_k(x; \mathcal{W})$
- We label now the multiclass targets via the 1-hot encoding, with the target of class  $k$  being the vector  $e_k = (0, \dots, \underbrace{1}_k, \dots, 0)$
- Then, if  $x \in C_k$ , its label is  $y = (y_1, \dots, y_K)^t$  is  $e_k$ , with  $y_k = 1$ , and

$$P(k|x) = P(k|x)^1 = P(k|x)^{y_k} = \prod_{c=1}^K P(c|x)^{y_c} \simeq \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c}$$

### The Sample Likelihood

- We have a sample  $S = \{x^p, y^p\}$  with  $(x^p, y^p)$  in class  $c(p)$  and, hence,  $y^p = e_{c(p)}$
- We also consider a posterior model  $P(c|x, w) = F_c(x, \mathcal{W})$ , the probability of getting  $S = (X, Y)$  is

$$\begin{aligned} P(Y|X; \mathcal{W}) &= \prod_1^N P(c(p)|x^p; \mathcal{W}) = \prod_1^N \prod_{c=1}^K P(c|x^p; \mathcal{W})^{y_c^p} \\ &\simeq \prod_1^N \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c^p} \end{aligned}$$

- Thus, we estimate  $\mathcal{W}$  by maximizing the sample's **likelihood**

$$\mathcal{P}(Y|X; \mathcal{W}) = \prod_{p=1}^N \prod_{c=1}^K F_c(x^p; \mathcal{W})^{y_c^p}$$

### The Cross Entropy Loss

- As before, we will work with the sample's log-likelihood, i.e.

$$\begin{aligned} \ell(\mathcal{W}; S) &= \log \mathcal{P}(Y|X; \mathcal{W}) = \sum_{p=1}^N \sum_{c=1}^K \log \left( F_c(x^p; \mathcal{W})^{y_c^p} \right) \\ &= \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W}) \end{aligned}$$

- The **cross-entropy** loss is now simply the negative log-likelihood  $-\ell(\mathcal{W}; S)$ , i.e.,

$$L(\mathcal{W}) = -\ell(\mathcal{W}; S) = - \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W})$$

- It is now straightforward to carry this into a NN setting


### Takeaways on Clasical MLPs

1. They have a layered structure with outputs computed in a **forward pass** using differentiable activations
2. Usual activations: ReLUs, sigmoid, tanh, linear
3. MLPs are **universal approximators**: this is indispensable for regression but has to be handled with care
4. MSE is the usual regression cost; we use a variant of cross entropy in classification
5. The error function gradients are computed by **backpropagation** of generalized errors
6. Backprop is basically a very simple procedure than can be **largely automated**
7. Regression and classification problems fit very well in this setting

## 2.3 MLP Regularization

### Overfitting in MLPs

- Since MLPs are a UAF, they can also approximate the noise in the sample
  - Given  $S = \{(x^p, y^p)\}$  if we allow enough hidden units in a SHL MLP we can arrive to a  $w^*$  s.t.  $y^p = f(x^p; w^*)$
  - We get thus a sample error  $\hat{e}(w^*) = 0$  but possibly with a very high generalization error
- I.e., MLPs may have very small bias but possibly large variance



(Ph.D. Thesis of Carlos Alafiz)

### Regularization vs Overfitting

- Why is there overfitting?
  - Because we may end up having too many weights with respect to sample size
  - Because we allow these weights to explore the entire weight space
- We can avoid this wandering if we limit  $w$ 's growth, for which we add a **regularization** term  $g(\|w\|)$  to  $e(w)$  with  $g(r)$  increasing
- Working with  $e_R(w) = e(w) + g(\|w\|)$  we have to **balance** the minimization of  $e(w)$  and that of  $g(\|w\|)$

- This balanced learning results in better generalization

### $L_2$ Regularization

- The simplest regularization procedure adds a quadratic penalty to the square error  $e$

$$e_R(w) = e(w) + \frac{\lambda}{2} \|w\|^2,$$

with  $\lambda$  the **weight decay** factor

- Also known as Tikhonov's regularization (or Ridge Regression for linear models)
- The desired effect is to constrain the evolution of  $w$ :
  - In fact, the regularized loss is essentially the **Lagrangian of the constrained problem**

$$\min_w e(w) \text{ subject to } \|w\|^2 \leq \rho, \quad \rho > 0$$

- The gradient becomes  $\nabla e_R(w) = \nabla e(w) + \lambda w$  and gradient descent becomes

$$w^{k+1} = w^k - \rho_k (\nabla_w e(w^k) + \lambda w^k)$$

- And the Hessian is  $\mathcal{H}_R(w) = \mathcal{H}(w) + \lambda I$

### How to choose $\lambda$

- Again, the correct choice of  $\lambda$  is crucial
- A small  $\lambda \ll 1$  results in a small regularization effect and overfitting risk appears
- A large  $\lambda \gg 1$  causes learning to forget about  $e(w)$  and the model will be essentially constant and will underfit
- Usually  $\lambda$  is chosen by using CV by exploring a discrete set of values  $\lambda_j$ ,
- The same can essentially be done for any other hyper-parameter:
  - Number of hidden layers and of hidden units?
  - Learning rate??
  - Minibatch size???

### MLP Ensembles

- Recall that for MLPs  $e(w)$  does not have a single minimum
- Moreover, the final MLP depends on the random initial  $w^0$
- And mini-batch training adds extra randomness to the final model
- Therefore, **you may never get the same MLP twice!!!**
- But we can turn this to our advantage by

- Starting from  $K$  independent initial weights and get  $K$  optimal weight sets  $w_k^*$
- Giving as the output the average  $f_e(x) = \frac{1}{K} \sum_1^K f(x; w_k^*)$
- We expect outputs of the form  $\hat{y}_k^p = y^p + \epsilon_k^p$  with the  $\epsilon_k^p$  independent
- Hence  $\frac{1}{K} \sum_k \epsilon_k^p \simeq 0$  and  $\frac{1}{K} \sum_k \hat{y}_k^p \simeq y^p$

### Takeaways on MLP Regularization

- MLPs have a **high risk of overfitting**
- Thus, they must be **regularized** to avoid overfitting
  - The regularization hyperparameter is chosen through **cross validation**
- MLP training has two **random components**: the initial point and minibatch selection
  - Thus, each training will converge to a **different optimum**
- MLP **ensembles** can take advantage of this
  - They imply extra costs but ensembles are **embarrassingly parallelizable**

## 2.4 Unconstrained Smooth Optimization


### Back to Optimization

- To build a ML model  $\equiv$  to **minimize a loss function**
- General optimization theory is a key tool in Machine Learning (ML)
- There are two optimization set ups in ML
  - **Unconstrained** optimization, slightly simpler and the one used for MLPs
  - **Constrained** optimization, wider and more complex
- In ML we have also to consider the optimization of differentiable and also non differentiable error functions
- MLP optimization: **unconstrained and differentiable**
- And also **batch**, i.e., over the entire sample, **mini-batch** over subsamples or **on line**, pattern by pattern

### Gradient Descent

- We “stretch” the weight set  $\mathcal{W}$  into a vector  $w$
- We recall that  $-\nabla e(w)$  is the maximum descent direction
- First idea: to build a (hopefully convergent) sequence  $w^k$  iterating (small) steps along  $-\nabla e(w^k)$

### Gradient Descent II



- In more detail, we start from a random  $w^0$  and compute

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k)$$

- $\rho_k$  is the **learning rate** (LR)
- With a small  $\rho_k$  we ensure  $e(w^{k+1}) < e(w^k)$  (although with possibly a very small descent)
- We can get a better iteration  $w^{k+1} = w^k - \rho_k^* \nabla_w e(w^k)$  using a  $\rho_k^*$  given by

$$\rho_k^* = \arg \min_{\rho} e(w^k - \rho \nabla_w e(w^k));$$

this is known as **line minimization**

- These GD methods are called **first order methods** in part because they only use  $\nabla e$

### One Dimensional Newton's Method

- Assume a quadratic function  $q(w) = aw^2 + bw + c$ , with  $a > 0$ , and a minimum at  $w^*$
- We can reach  $w^*$  from any  $w$  with a step  $\Delta w$  such that

$$0 = q'(w + \Delta w) = 2a(w + \Delta w) + b$$

- We have thus  $\Delta w = \frac{-b-2aw}{2a}$ , that is,

$$w^* = w - \frac{2aw + b}{2a} = w - \frac{1}{q''(w)} q'(w)$$

- This leads to **Newton's method**: minimize a general  $f$  iteratively using steps

$$w^{k+1} = w^k - \rho_k \frac{1}{f''(w^k)} f'(w^k)$$

with  $\rho_k$  a suitable learning rate

### One Dimensional Newton's Method II

- Notice how the  $\frac{1}{f''(w_k)}$  acts as a **self adjusting learning rate**
- When the parabola is very sharp (i.e.,  $a \gg 1$ ), we will be close to the minimum and the descent can overstep it
- But then  $f''(w_k)$  will be big and, hence,  $\frac{1}{f''(w_k)}$  **moderates** the gradient step
- Conversely, when the parabola is very wide (i.e.,  $a \simeq f''(w_k) \ll 1$ ), the minimum will be far away and the gradient step small descent will
- But then  $f''(w_k)$  will be rather small and, hence,  $\frac{1}{f''(w_k)}$  **accelerates** the gradient step
- We can extend this to several dimensions

### Multidimensional Newton's Method

- For a  $d$  dimensional  $w$ , the Taylor expansion of  $e$  at an optimum  $w^*$  is

$$e(w) \approx e(w^*) + \frac{1}{2}(w - w^*)^t \cdot \mathcal{H}(w^*) \cdot (w - w^*)$$

–  $\mathcal{H}(w^*)$  is the **Hessian** of  $e$  at  $w^*$  and  $\nabla e(w^*) = 0$

- It follows that  $\nabla e(w) \approx \mathcal{H}(w^*) \cdot (w - w^*)$  and, therefore,

$$w^* \approx w - \mathcal{H}(w^*)^{-1} \nabla_w e(w)$$

- This suggest to derive the  $w^k$  by

$$w^{k+1} = w^k - \rho_k \mathcal{H}(w^k)^{-1} \nabla_w e(w^k)$$

which is known as **Newton's Method** (NM)

### Variants of Newton's Method

- Theoretically NM converges very fast near  $w^*$ , but
  - Far from  $w^*$  convergence is not guaranteed
  - Moreover  $\mathcal{H}(w^k)$  may not be invertible
  - Besides, computing  $\mathcal{H}(w^k)$  is cumbersome and costly
- The **Gauss–Newton (GN)** approximation

$$\mathcal{H}(w) \simeq E[\nabla e(w) \nabla e(w)^T],$$

holds for any quadratic cost and simplifies the third problem

- The **Levenberg–Marquardt (LM)** method deals with the other two problems combining



- Gradient descent “away” from  $w^*$
- Gauss–Newton “near”  $w^*$

although “away” and “near” have to be properly addressed

### One-dimensional GN Approximation

- Assume  $e(w) = \frac{1}{2} \int (f(x; w) - y)^2 p(x, y) dx dy$ ; then

$$\begin{aligned} e'(w) &= \int (f(x; w) - y) \frac{\partial f}{\partial w}(x, y) p(x, y) dx dy; \\ e''(w) &= \int \left( \frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy \\ &\quad + \int (f(x; w) - y) \frac{\partial^2 f}{\partial w^2}(x, y) p(x, y) dx dy \end{aligned}$$

- Near a minimum  $w^*$  we may expect  $f(x; w) \simeq y$  and, therefore,

$$e''(w) \simeq \int \left( \frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy > 0$$

- Thus, for square errors, we can use first derivatives to approximate  $e''(w)$

### General GN Approximation I

- In the general case we have

$$\nabla e(w) = E [\nabla f(x; w) (f(x; w) - y)]$$

- And, therefore,

$$\begin{aligned} \nabla^2 e(w) &= E [\nabla^2 f(x; w) (f(x; w) - y)] + \\ &\quad E [\nabla f(x; w) \nabla f(x; w)^\tau] \end{aligned}$$

- The second term is easy to compute once we have  $\nabla f$
- If  $w \approx w^*$ ,  $f(x; w) \approx y$ ; therefore  $f(x; w) - y \approx 0$ ,
  - We can ignore the first, more complex, term

### General GN Approximation II

- We arrive at  $\nabla^2 e(w) \simeq E [\nabla f(x; w) \nabla f(x; w)^\tau]$  or, equivalently,

$$\mathcal{H}_{(i,j)(p,q)}(w) = \left( \frac{\partial^2 e}{\partial w_{ij} \partial w_{pq}}(w) \right) \simeq \left( E \left[ \frac{\partial f}{\partial w_{pq}} \frac{\partial f}{\partial w_{ij}} \right] \right)_{(i,j)(p,q)}$$

- $\mathcal{J} = E [\nabla f(x; w) \nabla f(x; w)^T]$  is **Fisher's information matrix**
  - $\mathcal{J}$  is **semidefinite positive**
  - And thus not necessarily invertible
- Often only its diagonal is considered and we have


$$\mathcal{H}_{(i,j)(i,j)}(w) = \left( \frac{\partial f}{\partial w_{ij}} \right)^2$$

### Advanced Optimization

- There are many more proposals in unconstrained optimization
- The **Conjugate Gradient** (CG) and **Quasi-Newton** (QN) methods are important in classical MLP training
- The basic idea in CG is to replace gradient descent directions  $g_k = -\nabla e(w_k)$  with new conjugate directions that try to keep somehow the previous “good directions”
- The basic idea in QN is to iterate as in NM but with simple approximations  $\mathcal{A}_k$  to  $\mathcal{H}^{-1}(w^k)$  that converge to  $\mathcal{H}^{-1}(w^*)$
- When training “small” NNs the **Limited-memory Broyden-Fletcher-Goldfarb-Shanno** (L-BFGS) QN variant is often used

### Accelerating Gradient Descent

- A possibility on the error surface of a NN is to have many long, narrow ravines
  - Gradients bounce in the narrow section, but may be very small along the wider one



### Momentum

- **Momentum** pushes them forward in the wider section (but we have to cope with a new parameter)
- Momentum tries to maintain descent's inertia with a term  $\Delta^k = w^k - w^{k-1}$ , i.e.,

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- More or less, what a **heavy ball** would do
- The goal is to keep  $w^k$  advancing in “plateaux”, i.e., small gradient zones
- Momentum can be seen as a crude approximation of a CG step
- Nice explanation at [Why Momentum Really Works](#)

### Nesterov’s Accelerated Gradient

- Let’s rewrite momentum in two steps
  1. Define  $\Delta^{k+1} = -\rho_k \nabla_w e(w^k) + \mu_k \Delta^k$  and
  2. Apply
 
$$w^{k+1} = w^k + \Delta^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- **Nesterov’s Accelerated Gradient** is a variant of this

$$\begin{aligned}\tilde{\Delta}^{k+1} &= -\rho_k \nabla_w e(w^k + \mu_k \tilde{\Delta}^k) + \mu_k \tilde{\Delta}^k; \\ w^{k+1} &= w^k + \tilde{\Delta}^{k+1}\end{aligned}$$

- In convex optimization it improves GD and is often highly effective in Deep Network training

### Adam

- Adam is currently the most widely used gradient-descent method for deep NN training
- At each step  $t$  Adam uses a new random mini-batch to
  - Update **exponentially smoothed** averages  $m_t$  of the gradient  $g_t$  and  $v_t$  of the the squared gradient  $g_t^2 = g_t \odot g_t$  as

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2;$$

- Since we have

$$E[m_t] \simeq (1 - \beta_1^t) E[g_t], \quad E[v_t] \simeq (1 - \beta_2^t) E[g_t^2],$$

we compute bias corrections  $\hat{m}_t, \hat{v}_t$  as

$$\hat{m}_t = \frac{1}{1 - \beta_1^t} m_t, \quad \hat{v}_t = \frac{1}{1 - \beta_2^t} v_t;$$

- Update weights as  $W_t = W_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$

### Understanding Adam

- Default values  $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999$ , and  $\epsilon = 10^{-8}$  usually work fine
- What is going on in Adam?
- At each step we work with estimates of the average local gradient:

$$\hat{m}_t \simeq E[\nabla_W f]; \quad \hat{v}_t \simeq \sigma(g_t) \simeq E[(\nabla_W f)^2]$$

- Thus, we can first see the Adam iterations

$$W_t = W_{t-1} - \alpha \frac{\widehat{m}_t}{\sqrt{\widehat{v}_t + \epsilon}}$$

as a kind of “normalized” gradient descent

- And in more detail, since we can expect  $E[(\nabla_W f)^2] \simeq E[\nabla_{W^2}^2 f]$ , we can see Adam as a “dampened” variant of diagonal Gauss–Newton’s steps

### When to Stop Training

- Typically the  $e(w_k)$  error diminishes towards an asymptotic minimum
  - If many units are used, we arrive to 0, which usually implies overfitting
- First solution: to use a separate **validation subset**  $V$  and stop training when the error in  $V$ , i.e., the **validation error** starts growing
  - But: How to choose  $V$ ? What do we do for small samples?
- Second solution (better): **get a good regularization** (and forget about overfitting)
  - A low CV error is also a low validation error
  - Now training stops because of reasons such as computational cost, but not because of overfitting risk

### Takeaways on MLP Optimization

- **Gradient descent** is the simplest method but possibly also the slowest
- **Momentum** can be used to speed it up
- **Newton’s method** is much faster but may be very costly and difficult to apply in full form
- All the previous methods require the (usually tricky) selection of a **learning rate**
- Second order methods such as Conjugate Gradient and **Quasi–Newton** avoid learning rates and are more efficient but costlier
- **Limited Broyden–Fletcher–Goldfarb–Shanno** (L-BFGS) is currently the option of choice for “small” MLPs

## 2.5 Computational Costs of MLPs

### MLPs and Big Data?

- Many Vs in Big Data: **Volume**, **Velocity**, Variety, Veracity, Value, ...
- Velocity: information flows in data streams that require fast processing and feed back
  - MLPs are rather fast
  - Less than linear models but much more so than SVMs

- Volume is probably the greatest attractive of Big Data
  - Huge samples and/or very large pattern dimension
  - Large impact in model training

### Volume in MLPs

- Parallelism is the first answer to Big Volume
- On a standalone machine it may be
  - Passive: let the SO distribute work among several CPU cores, or use GPUs with low level parallelized libraries such as Linpack or BLAS for linear algebra
  - Active: explicitly exploit a problem's parallelism programming an algorithm in, say, OpenMP
- It is easy to passively parallelize the training of several MLPs
  - But memory costs may multiply
- Training an isolated MLP is not CPU parallelizable:
  - There is a sequential layer dependence in Backprop
  - Each training iteration depends on the previous one
  - But matrix-vector multiplications can exploit GPU parallelism

### Handling Huge Sample Training

- Two consequences of a large  $N$  are
  - Sample doesn't fit in memory and we have to split it somehow
  - Training gets "lost", for in the global gradient  $\nabla e = E[\nabla e^\ell]$  we average many local gradients that may cancel each other out
- The first problem has been always present in fields such as analog signal filtering
- Solution for both: adaptive or **on line**, i.e., pattern by pattern, weight updates
  - It also allows a relatively simple setting for a theoretical analysis
  - But not used today; instead, medium-to-large NNs are trained using **mini-batches**

### MLP Complexity

- MLPs are fast to apply but costly to train
- How can we estimate this cost?
- MLP training cost is determined by
  - Its **architecture**, that determines the number of weights to fit and that is also dependent on the input dimension  $D$

- The full or mini-batch **sample size**  $N$ , that determines the cost of the averages to be computed
- The **training method**, with more or less iterations that, in turn, are more or less costly

### Forward Pass Complexity

- The number of weights in a single hidden layer (SHL) MLP with  $D$  inputs, one output and one hidden layer with  $H$  units is

$$(D + 1) \times H + (H + 1) \simeq HD$$

- Each extra hidden layer pair with  $(H_1, H_2)$  units adds  $(H_1 + 1) \times H_2 \simeq H_1 H_2$  weights
- The cost per pattern of the feedforward step between layers  $h - 1, h$  is basically that of  $H_{h-1} \times H_h$  floating point operations (FPOs)
- Thus, for a general MLP the cost in floating point operations of a forward pass is  $\simeq N \times (\sum_h H_h \times H_{h-1})$ 
  - Very fast on GPUs as it theoretically becomes  $O(N \times \sum_h H_h)$

### The Cost of Computing $\nabla e^\ell$

- Computing a local gradient  $\nabla e^\ell$  in a SHL MLP with square error and one output essentially requires to compute
  - $H$  components for the hidden to output connections, with a  $O(1)$  cost each, for  $\frac{\partial e^\ell}{\partial w_h^O} = (\hat{y} - y) o_h$
  - $DH$  components for the input to hidden connections with essentially an  $O(1)$  cost each, for  $\frac{\partial e^\ell}{\partial w_{h,j}^H} = \left( \sum_{i=1}^L \delta_i w_{ih}^O \right) \sigma'(a_h) x_j$
- In these we are omitting the cost of the partials with respect the bias, but they are much smaller
- Thus, the overall cost of computing  $\nabla e^\ell$  on a SHL MLP is essentially  $O(DH)$
- More layers with  $H_h$  units add a cost  $O(H_{h-1} \times H_h)$ 
  - $O(\sum_h H_h)$  on GPUs

### The Cost of Computing $\nabla e$

- For a mini-batch of size  $N_b$ , the cost of the mini-batch gradient  $\nabla e$  of a SHL MLP is  $O(N_b \times D \times H)$
- And extra layers add a cost  $O(N_b \times H_{h-1} \times H_h)$
- This is of the same order of magnitude than the cost of the forward pass
- This also dominates the  $O(H_{h-1} \times H_h)$  cost of updating the  $(H_{h-1}, H_h)$  weights in gradient descent

- And all these costs have to be multiplied by the number of training **epochs**

### Training Complexity

- The important term in the overall training cost is  $\text{nEps} \times \text{cost of } \nabla e$ , with nEps the number of epochs
  - One epoch = one pass on the entire sample
- Thus, the **globally dominant term** in a SHL MLP is

$$\text{nEps} \times \text{cost of } \nabla e = O(\text{nEps} \times N \times D \times H)$$

with  $N$  here the entire sample size

- And for each hidden layer we add an extra cost

$$O(\text{nEps} \times N \times H_{h-1} \times H_h)$$

- Thus, training many layered, large MLPs can be **very costly**
  - But **linear in sample size  $N$**

### Takeaways on MLP's Cost

- MLP complexity is determined by its **architecture**  $\{H_h\}$ , **training procedure** and **sample size**  $N$
- The forward and backward MLP passes have basically the **same complexity**
- Their cost per **pattern and layer** is  $H_{h-1} \times H_h$
- For gradient descent these costs are **multiplied by the number nEps of epochs**
- First order methods essentially **do not add extra complexity**
- Second order methods add extra per iteration costs but should **require less iterations**
- Single MLP training is **not CPU parallelizable**
- But **GPUs** can greatly improve MLP processing costs